

Sources

Every claim the panel-loop spec borrows, traced to the sentence it came from. Quotes are verbatim. The right-hand column says where each one is used in the design.

Andrej Karpathy — *Software Is Changing (Again)*

Y Combinator AI Startup School · June 2025 · full transcript PRIMARY

singjupost.com/andrej-karpathy-software-is-changing-again

written summary: latent.space/p/s3

"Usually they are doing the generation, and we as humans are doing the verification. **It is in our interest to make this loop go as fast as possible.**"

USED FOR

The whole premise – verification is the constraint, not generation quality.

"You are in charge of the autonomy slider... depending on the complexity of the task at hand, you can tune the amount of autonomy."

USED FOR

The opening slider diagram, and dialling autonomy per seat.

"I'm always scared to get way too big diffs. **I always go in small incremental chunks.**" — on "how we keep AI on the leash"

USED FOR

One decision per run. Batch size bounded by what you can check.

"A GUI utilizes your **computer vision GPU in all of our head. Reading text is effortful.**"

USED FOR

Step 4 – the visual verification surface.

"[A] GUI allows a human to **audit the work of these fallible systems and to go faster.**" · "It's much easier to just see a diff as, like, red and green change."

USED FOR

Before/after book pages with findings pinned, instead of a memo.

"[A] new category of consumer and manipulator of digital information. **It used to be just humans through GUIs or computers through APIs. And now we have a completely new thing... agents.**"

USED FOR

letters.json – the build manifest
panelists cite instead of inferring from CSS.

Andrew Ng — *Agentic Design Patterns, Part 1*

The Batch · 20 March 2024 PRIMARY

deeplearning.ai/the-batch/how-agents-can-improve-llm-performance

Reflection: "The LLM examines its own work to come up with ways to improve it."

USED FOR

The Critic role – your five seats.

Multi-agent collaboration: "More than one AI agent work together, splitting up tasks and discussing and debating ideas."

USED FOR

Round 2 – the rebuttal round.

HumanEval, coding tasks

GPT-3.5 zero-shot 48.1%

GPT-4 zero-shot 67.0%

GPT-3.5 + agentic workflow ... up to 95.1%

→ the workflow beat the model upgrade

Andrew Ng — letter on degrees of agency

The Batch, issue 253 · 12 June 2024 PRIMARY

deeplearning.ai/the-batch/issue-253

"Rather than arguing over which work to include or exclude as being a true agent, we can acknowledge that **there are different degrees to which systems can be agentic.**"

USED FOR

The autonomy dial – LOW / MEDIUM / HIGH per seat, rather than one setting for all five.

There is "a gray zone between what clearly is not an agent (prompting a model once) and what clearly is (say, an autonomous agent that, given high-level instructions, plans, uses tools, and carries out multiple, iterative steps of processing)."

USED FOR

LEARN and CHILD sit near the "prompt once" end on purpose.

Andrew Ng — *Evals and Error Analysis, Part 1*

The Batch · 15 October 2025 PRIMARY

deeplearning.ai/the-batch/improve-agentic-performance-with-evals-and-error-analysis-part-1

The "single biggest predictor of how rapidly a team makes progress building an AI agent" is its "ability to drive a disciplined process for evals... and error analysis."

USED FOR

Rung 3 – the decision log, and why it can't be skipped.

"Quickly build a prototype, then **manually examine a handful of agent outputs** to see where it performs well and where it stumbles" — then build "datasets and error metrics — sometimes objective metrics implemented in code, and sometimes subjective metrics using LLM-as-judge."

USED FOR

The eval timeline: log
→ count by hand → fix
→ only then LLM-as-judge.

Cognition — *Don't Build Multi-Agents*

Walden Yan, Cognition PRIMARY · the counter-argument this design has to survive
cognition.com/blog/dont-build-multi-agents

Principle 1. "Share context, and share full agent traces, not just individual messages."

USED FOR

Round 2 shares whole traces, not my summary of them.

Principle 2. "Actions carry implicit decisions, and conflicting decisions carry bad results."

USED FOR

The hard boundary: the panel reviews in parallel, the build is edited by one thread.

LangChain — *How and when to build multi-agent systems*

langchain.com/blog/how-and-when-to-build-multi-agent-systems

Read-heavy work parallelises cleanly; write-heavy work collides. Anthropic's own research system delegates the synthesis writing back to a single main agent.

USED FOR

Why a review panel is safe where parallel editing is not.

CORRECTION — 8 SEPT 2026

A number I got wrong the first time

An earlier draft said Ng's reflection split moved GPT-4 from roughly **67% to 88%** on HumanEval. That came from a secondary blog post and **does not appear in Ng's writing**. The figures above are his, from The Batch, 20 March 2024.

Two things to keep straight about them: the 95.1% is for *agentic workflows in general*, not for the Critic/Reviser split measured on its own — and it is a coding benchmark, not evidence about review panels. Cite it for the claim Ng actually makes: the workflow mattered more than the model upgrade.

NOT FROM ANY OF THEM

What this project invented

The five seats and their dials · the context packet and what it withholds · the triage split where verifiable findings skip the debate · the Actor / Critic / Reviser *naming* (Ng names the pattern **Reflection** and describes a separate critic; the three-role decomposition is the pattern literature's) · "confident garbage" as a phrase · the visual surface design · the five-rung ladder · `fr-lint.js` and every Fade Readers rule it checks.

These are argued from the sources above. They are not endorsed by them.